

# TSSL Lab 1 - Autogressive models

We load a few packages that are useful for solving this lab assignment.

```
In [1]: import pandas # Loading data / handling data frames
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression # used for solving linear regression problems
from sklearn.neural_network import MLPRegressor # used for NN model
from tskitools.lab1 import acf, acfplot # Module available in LISAN - used for plotting ACF
```

## 1.1 Loading, plotting and detrending data

In this lab we will build autoregressive models for a data set corresponding to the Global Mean Sea Level (GMSL) over the past few decades. The data is taken from <https://climate.nasa.gov/vital-signs/sea-level/> and is available on LISAN in the file `sea_level.csv`.


**Q1:** Load the data and plot the GMSL versus time. How many observations are there in total in this data set?

**Hint:** With pandas you can use the function `pandas.read_csv` to read the csv file into a data frame. Plotting the time series can be done using `pyplot`. Note that the sea level data is stored in the 'GMSL' column and the time when each data point was recorded is stored in the column 'Year'.

**A1:**

```
In [2]: sea_level = pandas.read_csv('sea_level.csv')

plt.plot(sea_level.Year, sea_level.GMSL)
plt.title('Sea Level vs Year')
plt.xlabel('Year')
plt.ylabel('GMSL')
```



In [3]: sea\_level.shape

Out[3]: (97, 2)

**Q2:** The data has a clear upward trend. Before fitting an AR model to this data need to remove this trend. Explain, using one or two sentences, why this is necessary.

**A2:** It is easier to work with stationary AR model, but in the above plot, we see that it is a non stationary model as it has an upward trend, the mean or expected value is changing with time. On removing the trend, we can make this as a stationary process.

**Q3:** Detrend the data using the steps:

1. Fit a straight line,  $\mu = \theta_0 + \theta_1 t$  to the data based on the method of least squares. Here,  $t_i$  is the time point when observation  $t$  was recorded.

**Hint:** You can use `lin.LinearRegression().fit(...)` from `sklearn`. Note that the inputs need to be passed as a 2D array.

Before going on to the next step, plot your fitted line and the data in one figure.

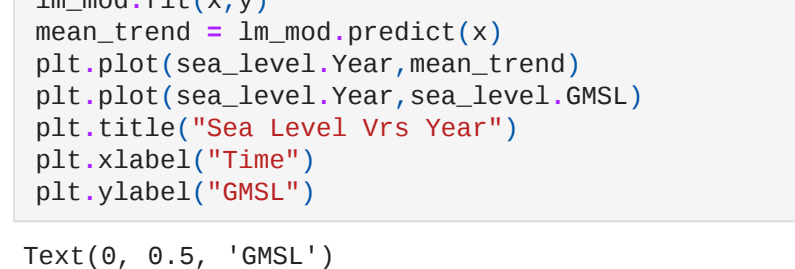
1. Subtract the fitted line from  $y_t$  for the whole data series and plot the deviations from the straight line.

**From now, we will use the detrended data in all parts of the lab.**


**Note:** The GMSL data is recorded at regular time intervals, so that  $t_{i+1} - t_i = \text{const.}$  Therefore, you can just as well use  $t$  directly in the linear regression function if you prefer,  $\mu = \theta_0 + \theta_1 t$ .

```
In [4]: # mu = np.array([1,1] for i in sea_level.Year.values)
# sea_level.Year.values
x = np.array(x.reshape(-1, 1))
y = sea_level.GMSL.values
```

```
In [5]: lin_mod = lin.LinearRegression()
lin_mod.fit(x,y)
mean_trend = lin_mod.predict(x)
plt.plot(sea_level.Year, mean_trend)
plt.plot(sea_level.Year, sea_level.GMSL)
plt.title('Sea Level vs Year')
plt.xlabel('Year')
plt.ylabel('GMSL')
```



```
In [6]: # flat_trend = np.array([y[i] - mean_trend[i] for i in range(len(y))])
# flat_trend = mean_trend
plt.plot(x, flat_trend)
plt.title('Detrended data')
plt.xlabel('Year')
plt.ylabel('Detrended GMSL')
```



**Q4:** Split the detrended time series into training and validation sets. Use the values from the beginning up to the 700th time point (i.e.  $y_t$  for  $t = 1$  to  $t = 700$ ) as your training data, and the rest of the values as your validation data. Plot the two data sets.


**Note:** In the above, we have allowed ourselves to use all the available data (train + validation) when detrending. An alternative would be to use only the training data also when detrending the model. The latter approach is more suitable if either:

- we view the linear detrending as part of the model choice. Perhaps we wish to compare different polynomial trend models, and evaluate their performance on the validation data, or
- we wish to use the second set of observations to estimate the performance of the final model on unseen data (in that case it is often referred to as "test data" instead of "validation data", in which case we should not use observations when fitting the model, including the detrending step.

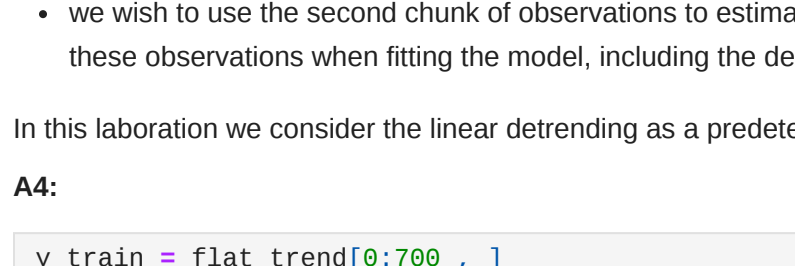
In this lab we consider the linear detrending as a predetermined preprocessing step and therefore allow ourselves to use the validation data when computing the linear trend.

**A4:**

```
In [7]: y_train = flat_trend[:700]
y_val = flat_trend[700:]
x_train = x[:700]
x_val = x[700:]
plt.plot(x_train, y_train)
plt.title('training data')
```



```
In [8]: plt.plot(x_val, y_val)
plt.title('Validation data')
```



## 1.2 Fit an autoregressive model

We will now fit an AR(p) model to the training data for a given value of the model order  $p$ .

**Q5:** Grab a function that fits an AR(p) model to an arbitrary value of  $p$ . Use this function to fit a model of order  $p = 10$  to the training data and write out (or plot) the coefficients.

**Hint:** Since fitting an AR model is essentially just a standard linear regression we can make use of `lin.LinearRegression().fit(...)` similarly to above. You may use the template below and simply fill in the missing code.

**A5:**

```
def fit_ar(p):
    """Fits an AR(p) model. The loss function is the sum of squared errors from t=1 to t=n.
    :param p: int, the training function
    :param p: int, AR model order
    :return theta: array (p,1), learnt AR coefficients
    """
    # Number of training data points
    n = len(y_train) # <COMPLETE THIS LINE>
    # Construct the regression matrix
    Phi = np.zeros(shape = (n,p,p)) # <COMPLETE THIS LINE>
    for j in range(p):
        Phi[:,j] = y_train[p-j-1:n-j-1] # <COMPLETE THIS LINE>
    # Drop the first p values from the target vector y
    yy = y_train[p:n] # yy = y[(n-p):], ..., y[n]
    # Here we use fit_intercept=False since we do not want to include an intercept term in the AR model
    regr = lin.LinearRegression(fit_intercept=False)
    regr.fit(Phi,yy)
    return regr.coef_
```

```
In [10]: #p = [1,2,3,4,5,6,7,8,9,10]
#np.transpose([2,3])
p = 10
theta = fit_ar(y_train,p)
print(theta)
```

Out[10]:

[ 0.6258602 0.3978277 0.1584657 0.1745793 -0.6218479 -0.6595486  
 0.9997898 0.0788523 0.1157939 0.6288308]

**Q6:** Next, write a function that computes the one-step-ahead prediction of your fitted model. One-step-ahead here means that in order to predict  $y_t$  at  $t = t_0$ , we use the actual values of  $y_t$  for  $t < t_0$  from the data. Use your function to compute the predictions for both training data and validation data. Plot the predictions together with the data (you can plot both training and validation data in the same figure). Also plot the residuals.

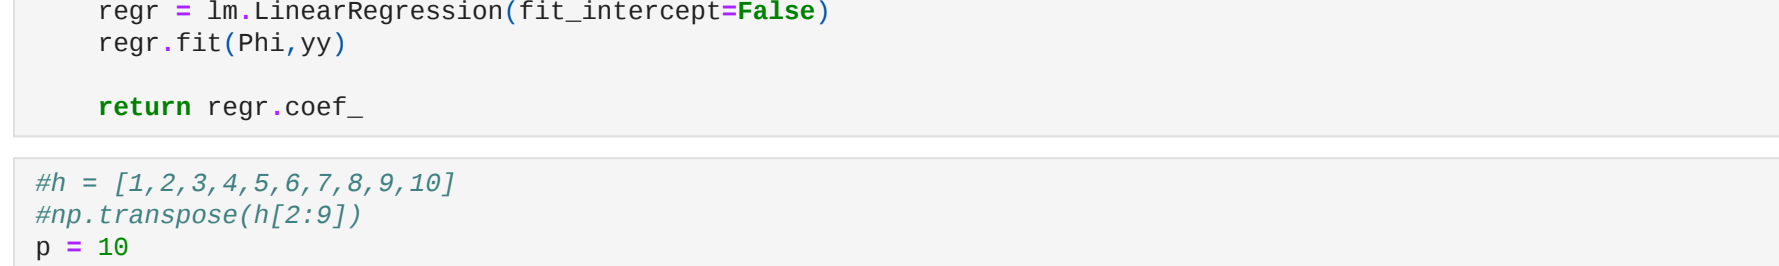
**Hint:** It is enough to call the predict function once, for both training and validation data at the same time.

**A6:**

```
def predict_ar_step(theta, y_target):
    """Predicts the value y_t for t = p+1, ..., n, for an AR(p) model, based on the data in y_train using
    one-step-ahead prediction
    :param theta: array (p,1), AR coefficients, theta(a1,a2,...,ap)
    :param y_train: array (n,1), the data points used to compute the predictions.
    :return y_pred: array (n-p), the one-step predictions (that y_train[p:], ..., that y_n)
    """
    n = len(y_train)
    p = len(theta)
    # Number of steps in prediction
    m = n-p
    y_pred = np.zeros(m)
    for i in range(m):
        # <COMPLETE THIS CODE BLOCK>
        y_pred[i] = np.dot(np.transpose(theta),Phi[i]) # <COMPLETE THIS LINE>
    return y_pred
```

```
In [12]: y_pred_whole = predict_ar_step(theta,flat_trend)
train_pred = y_pred_whole[:700-p]
val_pred = y_pred_whole[700-p:]

#residuals
resid_whole = flat_trend[p:] - y_pred_whole
plt.figure(figsize=(15,5))
plt.plot(flat_trend[:700],train_pred,'r-',label = "Train prediction")
plt.plot(flat_trend[700:],val_pred,'g-',label = "Validation prediction")
plt.plot(flat_trend[:700],label = "Time series")
plt.xlabel('t')
plt.ylabel('GMSL')
plt.legend()
```

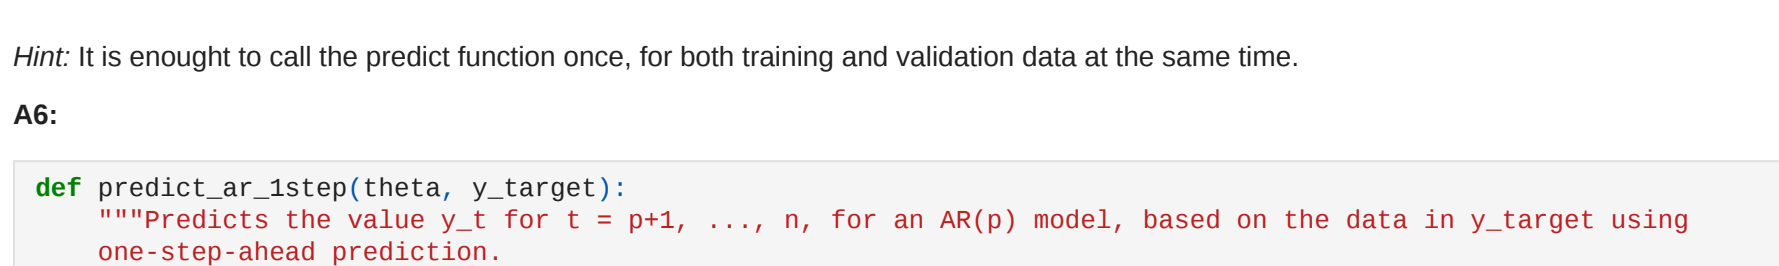


In [13]: #plotting residuals

fig, ax = plt.subplots(figsize = (15,5))

ax.plot(resid\_whole)

ax.set(title = "Residual plot", xlabel = "t")



**Q7:** Compute and plot the autocorrelation function (ACF) of the residuals only for the validation data. What conclusions can you draw from the ACF plot?

**Hint:** You can use the function `acfplot` from the `tskitools` module, available on the course web page.


**A7:**

In [14]: help(acfplot)

**Help:** on function acfplot in module tskitools.lab1:

```
acfplot(x, lags=None, conf=0.95)
Compute the empirical autocorrelation function.
:param x: array (n,1), sequence of data points
:param lags: int, maximum lag to compute the ACF for. If none, this is set to n-1. Default is None.
:param conf: float, number in the interval (0,1) which specifies the confidence level (based on a central limit theorem under a white noise assumption) for two dashed lines drawn in the plot. Default is 0.95.
:return: the autocorrelation function
```

```
In [15]: acfplot(resid_whole[690:])
```



Ideally we want residuals to be uncorrelated. That is to say it is good to have autocorrelation factor near zero for all lags (except 0).

## 1.3 Model validation and order selection

Above we set the model order  $p = 10$  quite arbitrarily. In this section we will try to find an appropriate order by validation.

**Q8:** Write a loop in which AR models of orders from  $p = 2$  to  $p = 150$  are fitted to the data above. Plot the training and validation mean-squared errors for the one-step-ahead predictions versus the model order.

Based on your results:

- What is the main difference between the changes in training error and validation error as the order increases?
- Based on these results, which model order would you suggest to use and why?

**Note:** There is no obvious "correct answer" to the second question, but you still need to pick an order to motivate your choice!

**A8:**

```
In [16]: from sklearn.metrics import mean_squared_error
error_train = []
error_val = []
for p in np.arange(2,151):
    theta = fit_ar(y_train,p)
    y_pred_whole = predict_ar_step(theta,flat_trend)
    train_pred = y_pred_whole[:700-p]
    val_pred = y_pred_whole[700-p:]
    error_train.append(mean_squared_error(y_train[p:],train_pred))
    error_val.append(mean_squared_error(y_val,val_pred))
```

```
#plotting
fig, ax = plt.subplots(figsize = (15,5))
ax.plot(p, error_train, 'r-', label = "Training Error")
ax.plot(p, error_val, 'b-', label = "Validation Error")
ax.set(title = "MSE", xlabel = "p", ylabel = "MSE")
ax.legend()
```

Out[16]:

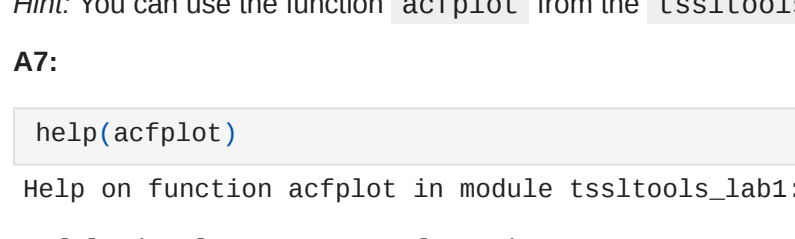
```
In [17]: error_val.index(min(error_val)) + 2 has p starts from 2
```

Out[17]: 70


a) We can see that the training error keeps decreasing as the order increases. But validation error decreases initially and then starts to increase as the order increases. b) If we have to choose a model, we feel it's better to pick the one with least validation error in the above case is the one with  $p = 70$ .

Based on the chosen model order, compute the residuals of the one-step-ahead predictions on the validation data. Plot the autocorrelation function of the residuals. What conclusions can you draw? Compare to the ACF plot generated above for  $p=10$ .

```
In [18]: p = 70
theta = fit_ar(y_train, p)
Phi = np.zeros(shape = (n,p,p)) # <COMPLETE THIS LINE>
y_pred_whole = predict_ar_step(theta,flat_trend)
y_train_pred = y_pred_whole[:700-p]
y_val_pred = y_pred_whole[700-p:]
acfplot(resid) # p is 70
```



In [19]: acfplot(resid\_whole[690:]) when p was 10



When compared to acfplot of  $p = 10$ , the new one (p = 70) has acf more closer to zero and within the CI. This is good, so we can say model with  $p = 70$  performs better than  $p = 10$

## 1.4 Long-range predictions

So far we have only considered one-step-ahead predictions. However, in many practical applications it is of interest to use the model to predict further into the future. For instance, for the sea level data studied in this lab, it is more interesting to predict the value one year from now, and not just 10 days ahead ( $10 \text{ days} = 1 \text{ time step}$  in this data).

**Q10:** With a function that simulates the value of an AR(p) model  $m$  steps into the future, conditionally on an initial sequence of data points. Specifically, given  $y_{1:n}$  with  $n \geq p$  the function code should predict the values

$$\hat{y}_{t:n} = \mathbb{E}[y_t | y_{1:n}], \quad t = n+1, \dots, n+m.$$

Use this to predict the values for the validation data ( $y_{700:97}$ ) conditionally on the training data ( $y_{1:700}$ ) and plot the result.

**Hint:** Use the pseudo-code derived at the first pen-and-paper session.

**A10:**

```
In [20]: def simulate_ar(y_train, theta, m):
    """Simulates an AR(p) model for m steps, with initial condition given by the last p values of y
    :param y_train: array (n,1), AR model parameters
    :param theta: array (p,1), AR coefficients, theta(a1,a2,...,ap)
    :param m: int, number of time steps to simulate the model for.
    """
    n = len(y_train)
    y_sim = np.zeros(m)
    Phi = np.transpose(theta)
    for i in range(m):
        y_sim[i] = np.dot(np.transpose(theta),Phi[i]) # <COMPLETE THIS LINE>
        Phi = Phi[1:n-p,:] # <COMPLETE THIS CODE BLOCK>
    return y_sim
```

```
In [21]: theta = fit_ar(y_train,70)
m = len(y_val)
p = len(theta)
y_sim = simulate_ar(y_train,theta,m)

fig, ax = plt.subplots(figsize = (15,5))
ax.plot(y_sim, 'c-', label = "Simulated")
ax.plot(y_val, 'b-', label = "True value")
ax.legend()
```

Out[21]:

**Q11:** Using the same function as above, try to simulate the process for a large number of time steps (say,  $m = 2000$ ). You should see that the predicted values eventually converge to a constant prediction of zero. Is this something that you would expect to see in general? Explain the result.

**A11:**

```
In [22]: m = 2000
y_sim = simulate_ar(y_train,theta,m)
fig, ax = plt.subplots(figsize = (15,5))
ax.plot(y_sim, 'c-', label = "Simulated")
ax.set(title = "Large simulation")
ax.legend()
```



As this is a stationary AR process (we have removed the upward trend) on simulating this to large number, the seasonal component reduces and goes to zero, as our theta is from -1 to 1 (dot product ends up near zero). (not really sure about this)

## 1.5 Nonlinear AR model

In this part, we switch to a nonlinear autoregressive (NAR) model, which is based on a feedforward neural network. This means that in this model the recursive equation for making predictions is still in the form

$\hat{y}_t = f(y_{t-1}, \dots, y_{t-p})$ , but this time  $f$  is a nonlinear function learned by the neural network. Fortunately almost all of the work for implementing the neural network and training is handled by the `sklearn` package with a few lines of code, and we just need to choose the right structure, and prepare the input-output data.

**Q12:** Construct a NAR(p) model with a feedforward (MLP) network, by using the `MLPRegressor` class from `sklearn`. Set  $p$  to the same value as you chose for the linear AR model above. Initially, you can use an MLP with a single hidden layer consisting of 10 hidden neurons. Train it using the same training data as above and plot the one-step-ahead predictions as well as the residuals, on both the training and validation data.

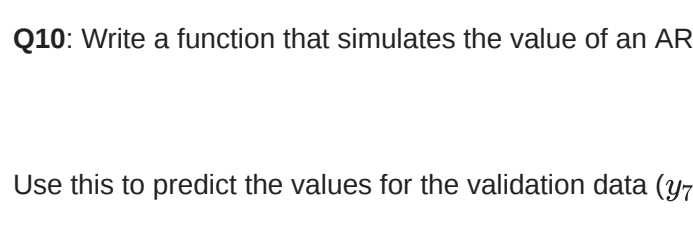
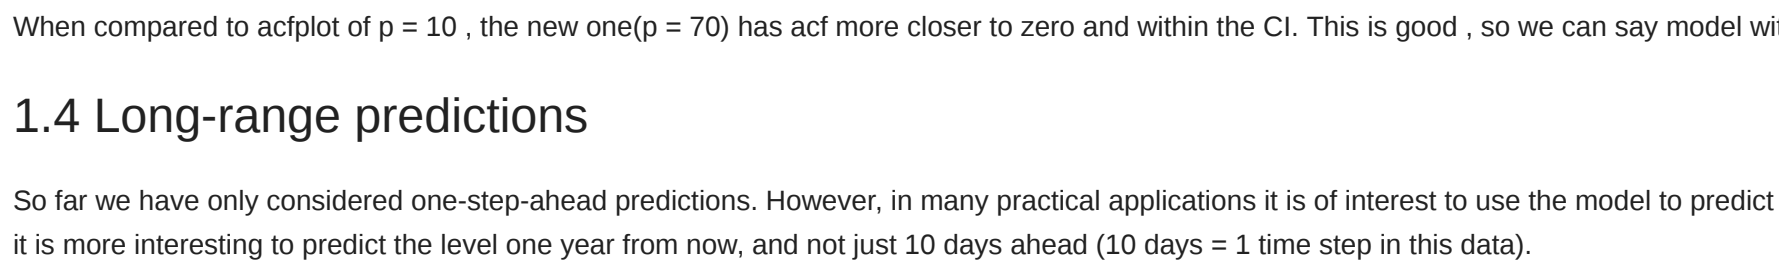
**Hint:** You will need the methods `fit` and `predict` of `MLPRegressor`. Read the user guide of `sklearn` for more details. Recall that a NAR model is conceptually very similar to an AR model, so you can reuse part of the code from above.

**A12:**

```
In [23]: p = 80
#copy code from fit_ar
# Number of training data points
n = len(y_train) # <COMPLETE THIS LINE>
# Construct the regression matrix
Phi = np.zeros(shape = (n,p,p)) # <COMPLETE THIS LINE>
for j in range(p):
    Phi[:,j] = y_train[p-j-1:n-j-1] # <COMPLETE THIS LINE>
# Drop the first p values from the target vector y
yy = y_train[p:n] # yy = y[(n-p):], ..., y[n]
# Using MLPRegressor instead of Linear Regression
mlp_model = MLPRegressor(hidden_layer_sizes = (10,), activation = 'relu', shuffle = False, max_iter = 5000)
mlp_model.fit(Phi,yy)
# Predicting
# To use predict, we need input of shape (n_samples,n_features)
n_samples = len(flat_trend)
ip_x = np.zeros(shape = (n_samples,p,p)) #inputs are the regressors
for i in range(p):
    ip_x[:,i] = flat_trend[p-i-1:n-p-i-1]
y_pred_whole = mlp_model.predict(ip_x)
train_pred = y_pred_whole[:700-p]
val_pred = y_pred_whole[700-p:]
resid_whole = flat_trend[p:] - y_pred_whole
```

```
f = plt.figure()
f.set_figsize(15)
f.set_figheight(5)
plt.plot(flat_trend,label = "Detrended data")
plt.plot(train_pred,'r-',label = "predicted training data")
plt.plot(val_pred,'g-',label = "predicted validation data")
plt.plot(flat_trend[:700],label = "Time series")
plt.xlabel('t')
plt.ylabel('GMSL')
plt.legend()
```

Out[23]:



**Q13:** Try to experiment with different choices for the hyperparameters of the network (e.g. number of hidden layers and units per layer, activation function, etc.) and the optimizer (e.g. `solvers` and `max_iter`).



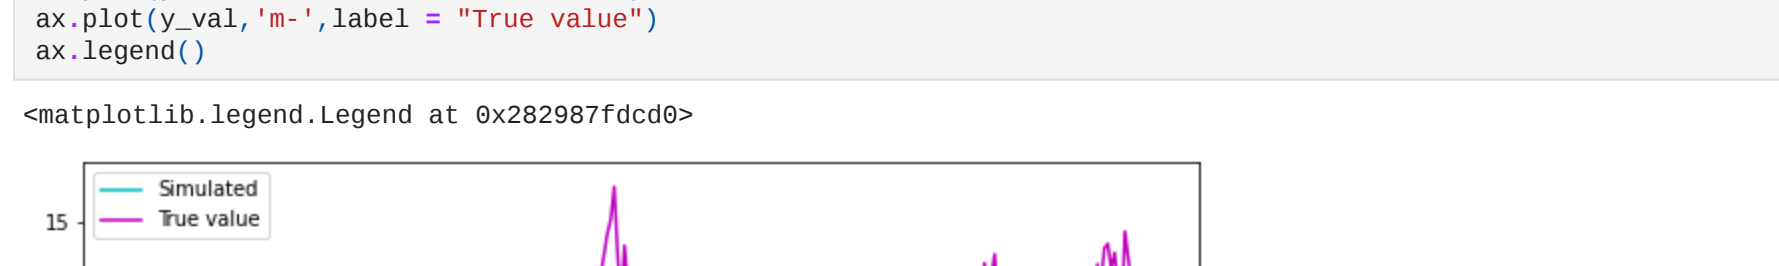
Are you satisfied with the results? Why/why not? Discuss what the limitations of this approach might be.

**A13:**

```
In [26]: p = 80
#copy code from fit_ar
# Number of training data points
n = len(y_train) # <COMPLETE THIS LINE>
# Construct the regression matrix
Phi = np.zeros(shape = (n,p,p)) # <COMPLETE THIS LINE>
for j in range(p):
    Phi[:,j] = y_train[p-j-1:n-j-1] # <COMPLETE THIS LINE>
# Drop the first p values from the target vector y
yy = y_train[p:n] # yy = y[(n-p):], ..., y[n]
# Using MLPRegressor instead of Linear Regression
mlp_model = MLPRegressor(hidden_layer_sizes = (10,10), activation = 'tanh', shuffle = False, max_iter = 1000, solver = 'adam')
mlp_model.fit(Phi,yy)
y_pred_whole = mlp_model.predict(ip_x)
train_pred = y_pred_whole[:700-p]
val_pred = y_pred_whole[700-p:]
resid_whole = flat_trend[p:] - y_pred_whole
```

```
f = plt.figure()
f.set_figsize(15)
f.set_figheight(5)
plt.plot(flat_trend,label = "Detrended data")
plt.plot(train_pred,'r-',label = "predicted training data")
plt.plot(val_pred,'g-',label = "predicted validation data")
plt.plot(flat_trend[:700],label = "Time series")
plt.xlabel('t')
plt.ylabel('GMSL')
plt.legend()
```

Out[26]:



We have kept two layer network with 10 hidden units, and 1000 max iter., increasing the iteration further did not give us any better results. Also adding more hidden units made the model worse due to probably overfitting. Acfplot has more values in the CI band while using adam compared to other options such as fmin or sgd.

In [ ]: